

HOW TO do good systems research and publish papers

Changwoo Min
ECE @ Virginia Tech
Oct 23rd, 2018

Executive Summary

- No magic wand
 - Read papers; think critically; build systems; measure performance; write papers
- But there are **anti-patterns** and **best practices**
- Following are some thoughts and suggestions that I learned from my advisers, collaborators, and colleagues.

Three most important things in doing systems research

- Problem selection
- Problem selection
- Problem selection

How to find a good problem

- What is a good problem?
 - Important problem
 - Unsolved problem (or problem that needs more solutions)
 - New problem
- What should I do?
 - Read papers; think critically; build systems; measure performance; write papers
- Where should I take a look?
 - Emerging new technologies (manycore, accelerator, RDMA, non-volatile memory)
 - Boundaries between layers (SSD vs. file system)
 - New aspects of a system (security, energy consumption)
 - New applications (IoT, Mobile, Cloud)
 - Interactions between subfields (OS, concurrency, database)

Why finding a good problem is not easy as a gradstudent?

- **You should change your role in learning**

- Undergrad: knowledge *consumer*
- Grad: knowledge *producer*

- **Are you ready to find an important problem?**

- Student: I have a cool idea, which is ...
- Adviser: Sounds great! The similar idea was published at 197x by someone who won the Turing award.

/ a month later */*

- Student: I have another cool idea, which is ...
- Adviser: Sounds great! The similar idea was published at 198x by someone who won the Turing award.

/ a month later */*

... 199x ...

/ a month later */*

... 200x ...`

Develop a good intuition

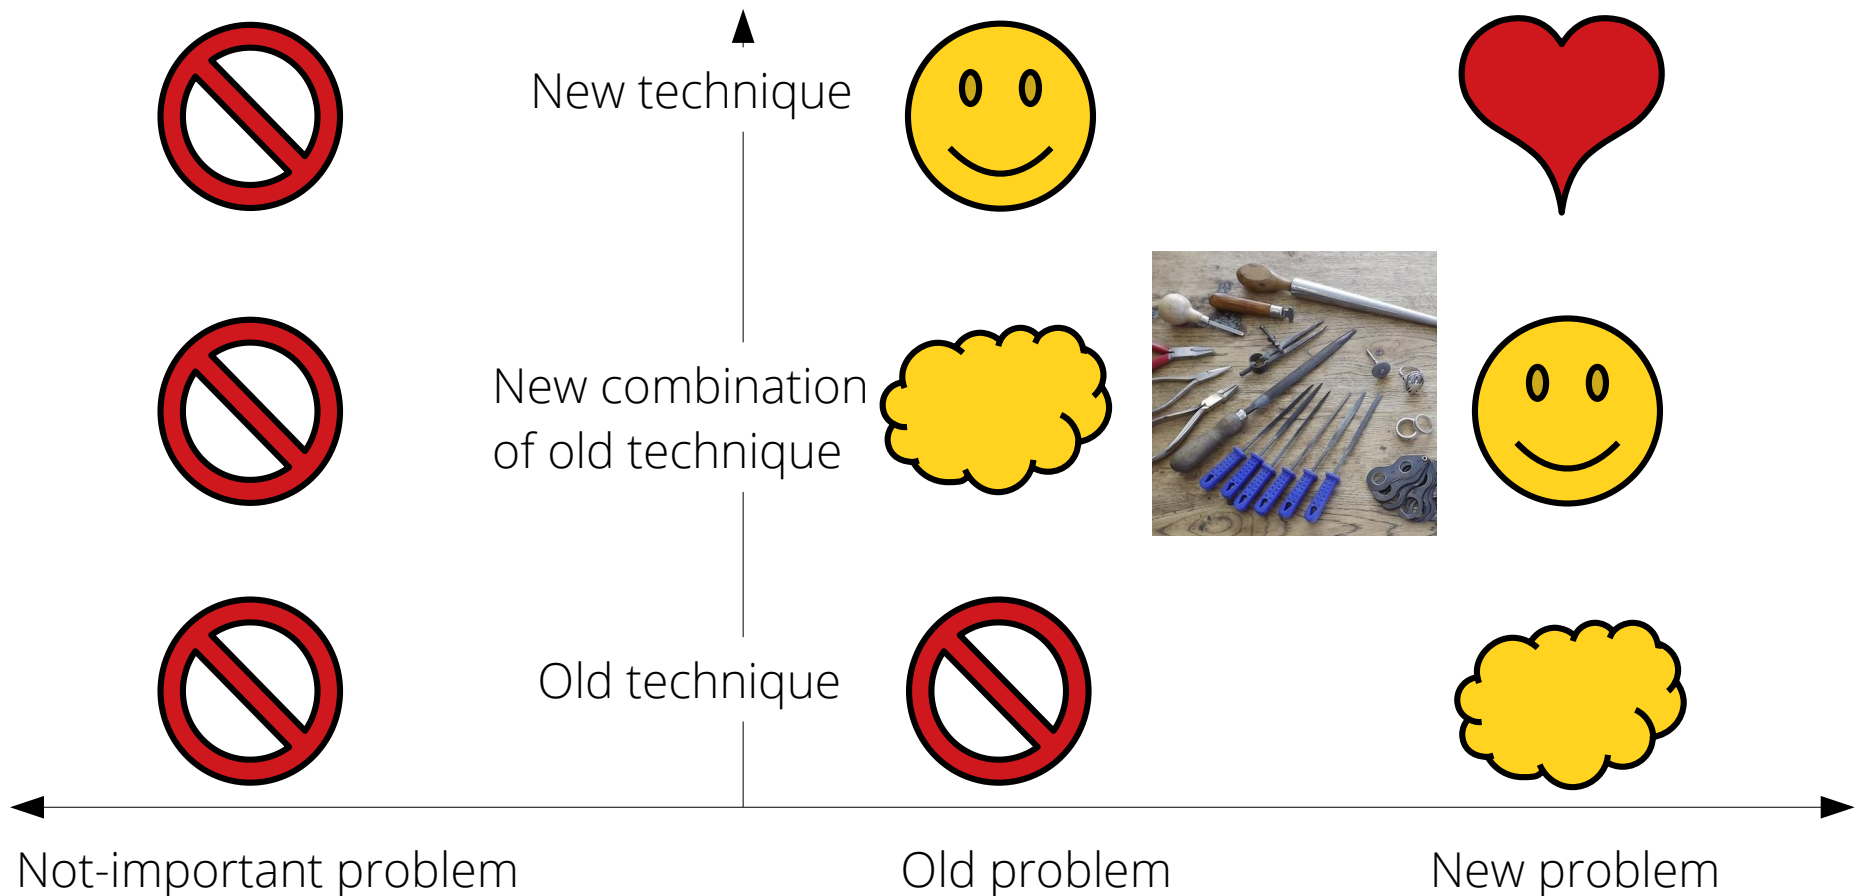
- How to approach the problem
 - What's the *fundamental* limitation of existing approaches?
 - Why do you think your approach would be better?
 - E.g., MV-RLU: two threads accessing different memory locations should run in parallel.
- Elevator pitch to your adviser, colleagues, senior meet in conferences
 - What are their first questions?
 - Why they are not enthusiastic to your idea?
- Improve your intuition/idea/algorithm
- Repeat
- Do you still think your approach is good?

Is your idea publishable? (1/2)

- Is your system better?
 - Better performance
 - Short tail latency
 - Save energy
 - Enable new types of applications
 - Secure by design
 - Etc.
- How much better is “better”?
 - 100x, 10x
 - 30%
 - ~~10%, 5%~~

Is your idea publishable? (2/2)

- Is your idea novel?



My rule of thumb to find an important problem and a good solution

: 30-20 rule

- First read 30 papers
 - Understand the problem space more deeply
 - Know the details of common techniques and the big ideas behind them
 - Know how to evaluate and what to compare
 - Find out the fundamental limitations of the 30 papers
 - Come up with your tentative solution to overcome the fundamental limitations
- Then read another 20 papers
 - Understand techniques proposed in 20 more papers
 - Test if your idea is competitive over the 20 papers at least one or two aspects (e.g., performance, energy, programmability, etc)

Is it doable?

- Can you develop the first prototype within a semester?
 - If no, ...
 - Learn necessary skills
 - Find good team mates
 - Limit the scope → Is your idea still competitive?
- Know your competition
 - What is the baseline?
 - What is the state-of-the-art?

Design / Implementation / Evaluation

- This is an iterative process.
 - Start the minimal design and implementation
 - Gradually add features
 - Or periodically rewrite the code
- Plan evaluation ahead
 - Write unit testcases
 - Integrate your benchmark and automate evaluation

Papers in systems research

- Romantic comedies



Papers in security research

- Superhero



Superhero Movies Treat Wo...
mic.com



Superhero Movie (2008) - I...
imdb.com



Top 10 Superhero Movies!
thesefantasticworlds.com



41 best Superhero Movies im...
pinterest.com



Top 50 Rotten Tomatoes S...
indiewire.com



Iron Man (2008 film) - Wiki...
en.wikipedia.org



Superhero Movies Treat Wo...
mic.com



Ant-Man (film) | Marvel Mov...
marvel-movies.wikia.com




Movie Review: "Captain A...
pinterest.com



Top 21 Superhero Movies ...
saltypopcorn.com.au



Superhero Movie Poster | ...
movieposteraddict.com



movie poster warehouse mov...
movieposter.com

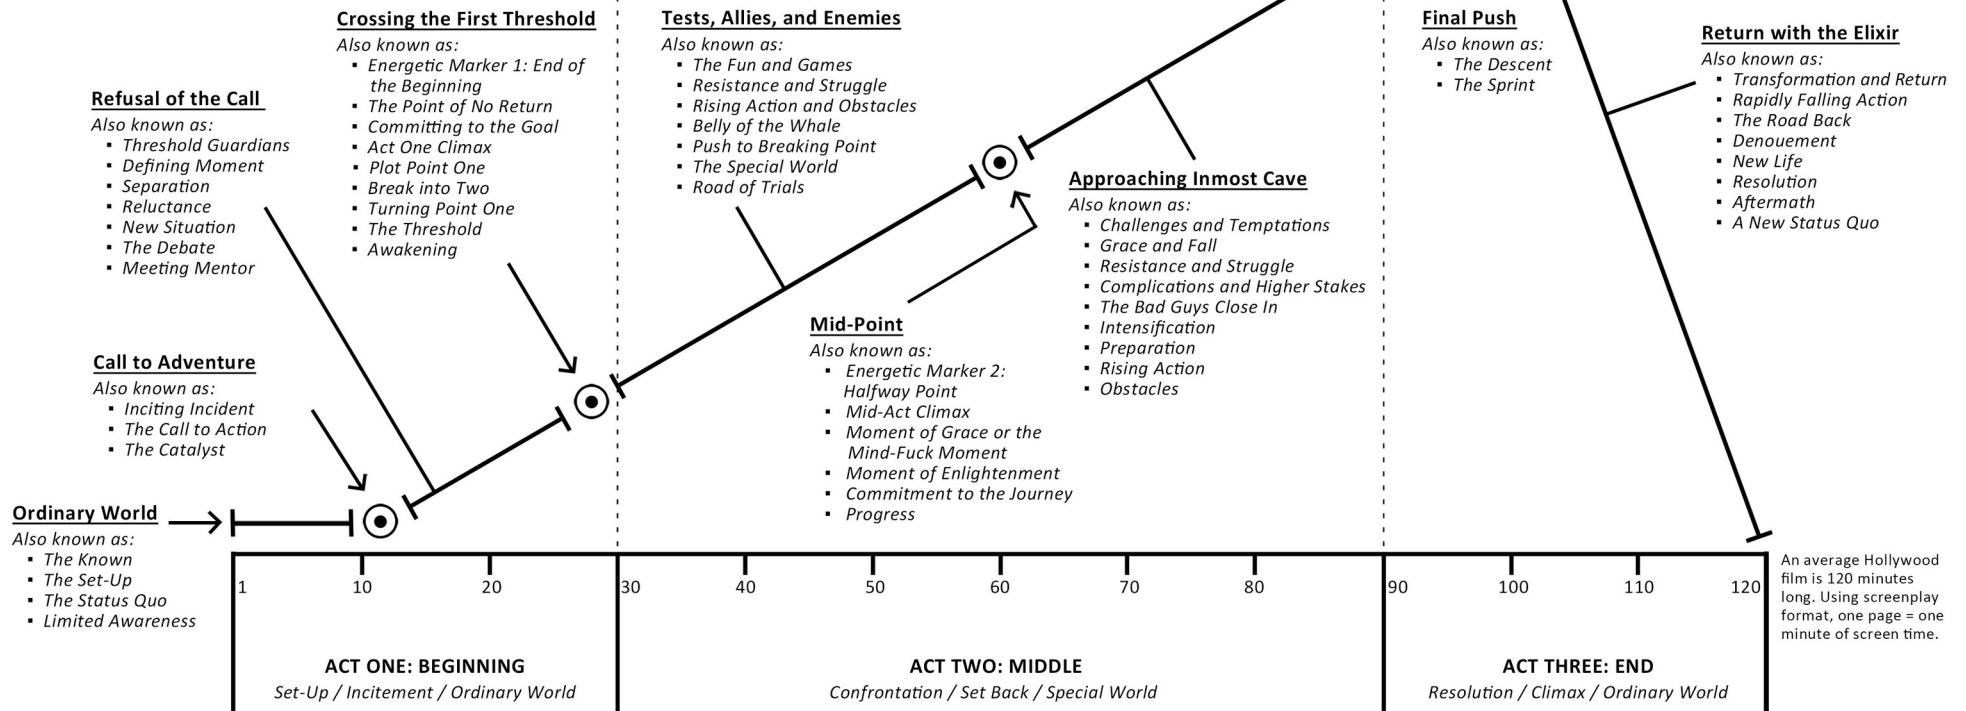
The Law of Genre

ARCHPLOT STRUCTURE

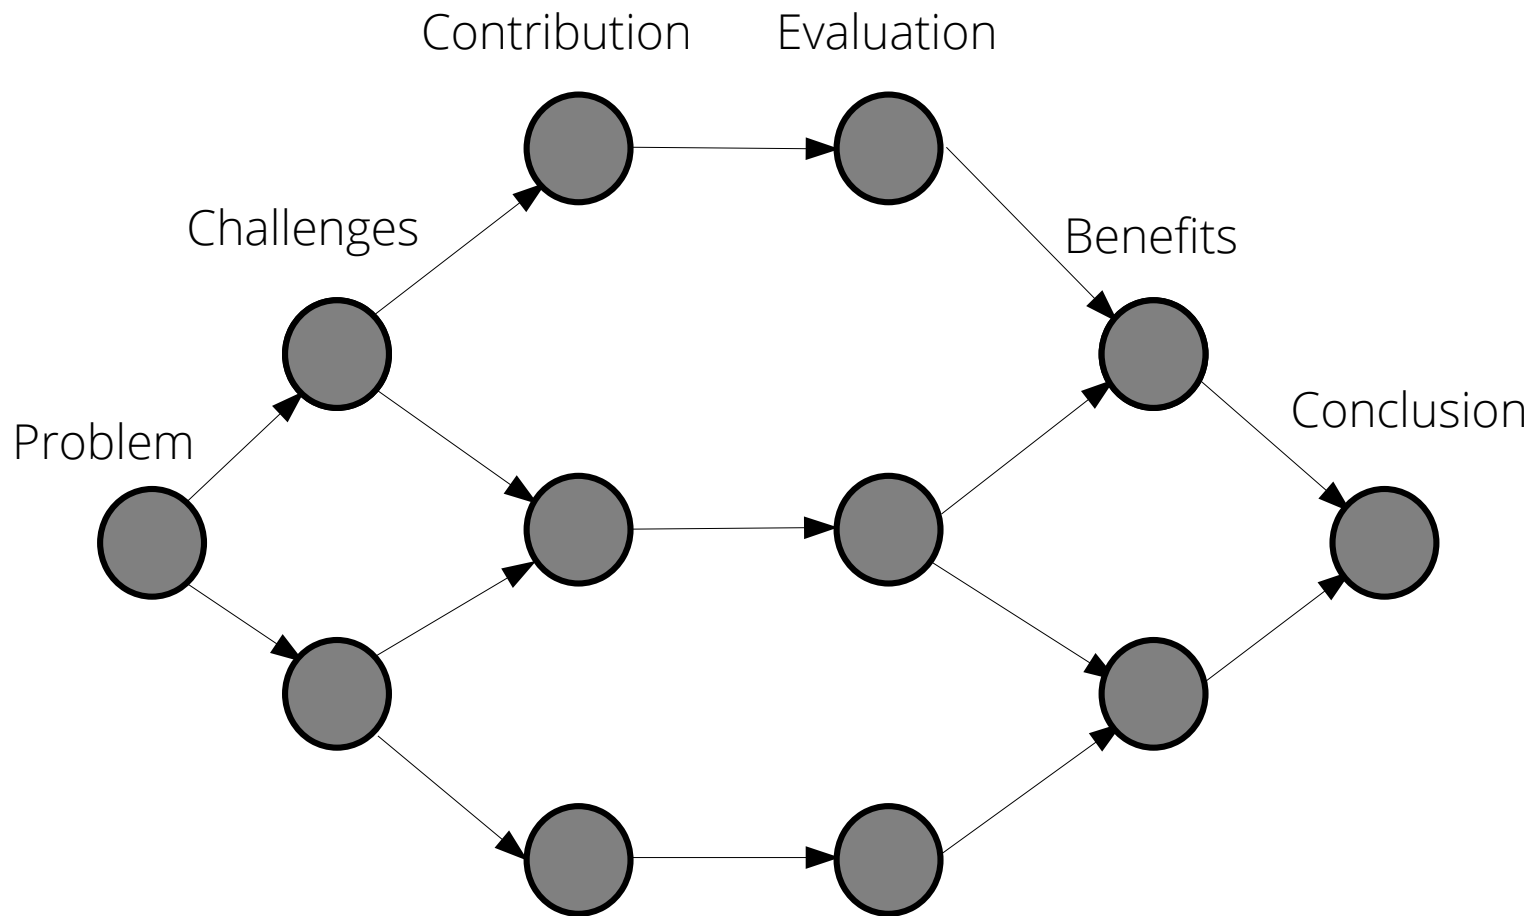
(AKA: Classic plot, the hero's journey, goal-oriented plot, Aristotelian story shape, *energeia* plot, and Hollywood screenwriting structure)

Archplot is a goal-oriented plot where, "for better or worse, an event throws a character's life out of balance, arousing in him the conscious and/or unconscious desire for that which he feels will restore balance, launching him on a Quest for his Object of Desire against forces of antagonism (inner, personal, extra-personal). He may or may not achieve it" (McKee, 196). **Film Examples:** *Toy Story*, *The Godfather*, *Back to the Future*, *Star Wars*, Etc. (Most American Hollywood films use arch plot). **Book Examples:** *Harry Potter* (Rowling), *Hunger Games* (Collins), *Speak* (Anderson), *Pride & Prejudice* (Austen), *Hamlet* (Shakespeare), *The Odyssey* (Homer), etc.

"I took a master class with Billy Wilder once and he said that in the first act of a story you put your character up in a tree and the second act you set the tree on fire and then in the third you get him down." - Gary Kurtz (Film Producer)



Archplot of a systems paper



Before writing a paper

- Know your strength and weakness
 - Read all related papers: 30 – 50
 - What is the real strength of my approach?
 - How will you defend your weakness?
- Does your evaluation results support your claim?
 - Get baseline numbers first
 - Run core evaluation results
- Is your system better? Really?

Typical outline of a systems paper

- Introduction
- Background and Motivation
- Design
- Implementation
- Evaluation
- Discussion and Limitation
- Related work
- Conclusion

Outlining your paper

- Top-down approach
- Find 2-3 most similar paper in style
 - Copy the style of the paper!
- Quickly outline the paper (within one day)
 - Decide section title
 - For each section
 - Key message of each paragraph
 - Evaluation
 - What to show? What to compare?
- Present the first version of outline to co-authors
 - What's their questions?
- Rewrite the outline from scratch (withing five hours)
 - Don't try to modify the first version
- Revise the outline
 - By co-authors

Outline examples

```
# MV-RLU: Scaling Read-Log-Update using Multi-Version Concurrency Control
```

Introduction

```
- **[[Set context]]**
```

- CDS is everywhere (e.g., kernel, database system, key-value store, DNS,
- *[CDS is the key to scale many software in manycore systems.]*

```
- **[[Why this is important?]]**
```

- More core -> small fraction of sequential execution -> performance coll
- Amdhals' Limit: (e.g., 1% sequential part)
 - 10 cores -> 9.2x
 - 32 cores -> 24.4x
 - 64 cores -> 39.2x
 - 128 cores -> 56.4x
 - 224 cores -> 69.3x
 - 448 cores -> 81.9x
- Such big machines are real
 - 64 cores -> a typical server in a data center (ref: Facebook)
 - 192 cores (24 core * 4 smt * 2 socket) -> new arm server (ref: news article)
 - 448 cores (28 core * 2 smt * 8 socket) -> Intel Xeon server (ref: r
- *[This is an increasingly important and interesting problem.]*

```
- **[[Introduce related work very briefly]]**
```

- Significant research efforts since the first inception of multiprocessors (refer Table 1 and Fig 1)
- Locking: mutual exclusion -> sequential execution
 - more parallelism while guaranteeing mutual exclusion
 - fine grained locking: linux VFS (XX -> YY)
 - more permissive lock:
 - inode mutex -> inode rswm: multiple readers
 - sequence lock in inode cache: multiple readers, single writer,
 - pros: intuitive, composable
 - cons: deadlock, livelock, complex, finer grained -> higher synch ov
- Lock free algorithm: memory hotspot, high synch cost
 - Difficult to design, hard to correct, not composable
 - Difficult to scalee
 - Scalable memory reclamation scheme

```
-UU-:~@*-F20 outline-v0.md Top (1,0) Git-master (Markdown Fly Helm A
```

1. Intro

- Tech trends
- Exim
- Summary of our finding
- Summary of contrigution
- Paper organization

2. Case study

- Multi-threaded IO is popular.
 - desktop [file is not a file]
 - mobile [revisiting, quasio]
 - cloud [rocks db, wiredtiger, ...]
- How applications make an effort to parallelisze file system operations?
 - And their results on ramdisk/ssd/hdd
- Exim
 - message file creation at spool dir. and move the the user dir.
 - use three processors
- RocksDB
 - Write performance in write-optimized LSM-like key/value store is limited by compaction performance.
 - Parallel compaction
 - merge two files into another two
- MariaDB/InnoDB
 - reader thread
 - writer thread
 - depending on operation (read/write) queue specific operation into separ
 - foreground-thread can continue work after metadata-operations done, doe
 - directory per database, file per table to allow parallel processing of

3. FxMark Benchmark Suite

- Goal
 - Identify hot spots in file system operations
 - Measure how identified hot spots affects real applications

- Micro benchmarks

- 18 benchmarks

- Application Benchmarks

- filebench: fileserver, varmail, oltp workloads
- Exim (mosbench), rocksdb (db bench), MariaDB (tpc-c)

```
-UU-:~@---F21 NOTE.paper Top (39,0) Git-master (Fundamental Helm A
```

Fix evaluation scenarios

- Decide graphs and tables in the evaluation section
 - X-axes, y-axes
- Get baseline number first
- Then get your numbers
- Automate everything

```
- [[Concurrent data structures]]
- [[Figure: 4 x 3 - double column]]
- Data structures
- linked list, hash table, binary tree (bst and citrus?)
- Size
- ??? 10K ???
- Read-write ratio
- read-only (0%)
- read-mostly (2%)
- read-intensive (20%)
- write-intensive (80%)
- Access skew: uniform random
- Comparisons
- RCU, RLU-gclk, MVRLU-ordo
- Harris-HP, Harris-QSBR,
- Citrus tree
- ??? SwissTM, STO ???
- ??? versioned programming, optik ???
- ??? Check synchrobench for easy comparison (lock-free
  algorithm?) ???
- [[Figure: comparison of abort ratio - 3 x 1 - single column]]
- x-axis: # of core
- y-axis: abort ration in %
- RLU, MVRLU, ??? STM (swisstm or sto)
- hash table: 2%-20%-80% (??? is hash table the best for this???)
- Explain why MV-RLU is good and why others show performance
  collapse
- analyze abort ratio, cache miss, perf prof, etc.
- [[Data set size]]
- [[Figure: 1 x 3 - single column]]
- Size: 1K, 10K, 100K
- Access skew: uniform random
- Comparisons
- RCU, RLU-gclk, MVRLU-ordo
- Harris-HP, Harris-QSBR,
- Citrus tree
- ??? SwissTM, STO ???
- ??? versioned programming, optik ???
```

Writing in bottom-up order

- Abstract
 - Introduction
 - Background and Motivation
 - Design
 - Implementation
 - Evaluation
 - Discussion and Limitation
 - Related work
 - Conclusion
- Evaluation
 - Design
 - Implementation
 - Background and Motivation
 - Related work
 - Discussion and Limitation
 - Introduction
 - Conclusion
 - Abstract

Writing a section

- Design the logical flow between paragraph
- Writing a paragraph
 - Conclusion first
 - Why? Provide supporting data

Evaluation Section

- ***Show you achieved all you promised***
 - ***Benefit? Design decision?***
- List up evaluation questions
 - Show your evaluation is complete
- Show benefit of your system first
 - then explain how your design decisions affects
 - X: micro-benchmark → application benchmark
 - O: application benchmark → micro-benchmark
- Draw graphs and tables

Design Section

- *Make sure non-expert reviewers can easily understand*
- Draw a self-contained, illustrative example of your system
 - In most cases, layered diagram is useless
 - Introduce key components and their interactions of your system
 - Define terms in your system
- Explain the diagram
- Emphasize your new, cool stuff

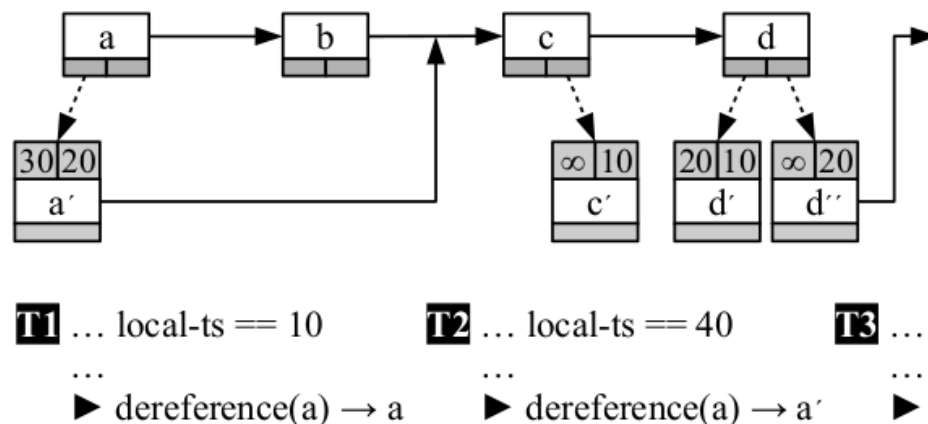
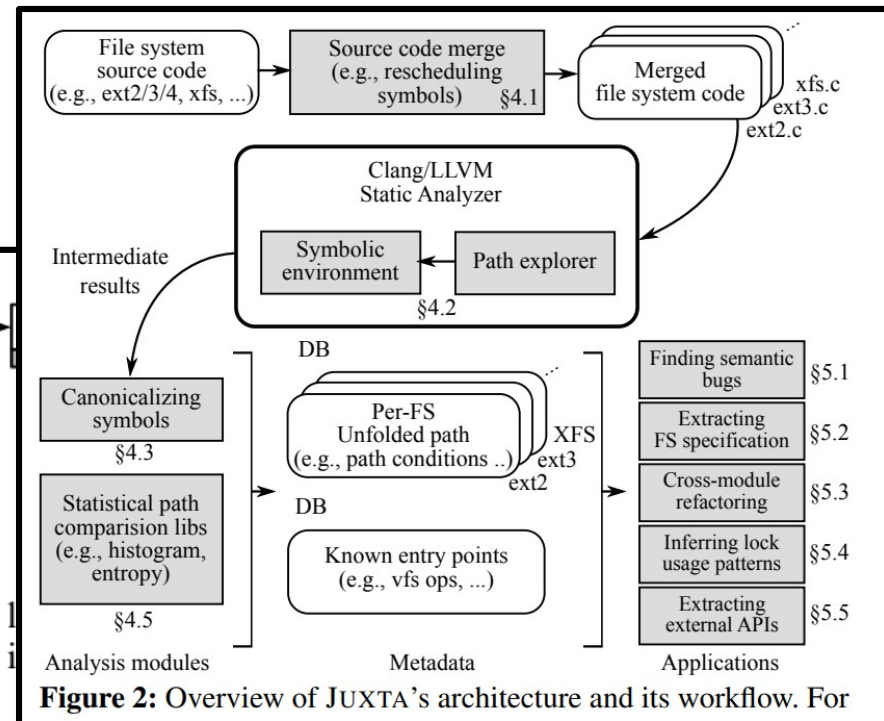


Figure 3: Illustrative snapshot of concurrent operations in the MV-RLU-based linked list. ▶ denotes where a thread executes. At time 30,



Implementation Section

- ***Clearly separate your reusable idea and implementation artifacts***
- Design section
 - reusable idea which will be usable in 20 years later
- Implementation section
 - Implementation-specific optimization
 - Show off your implementation is solid
 - Limitation in implementation
 - != limitation in design

Background and Motivation

- Background

- *Provide enough background for non-expert reviewers to understand this paper without reading other paper*

- Motivation

- *Show fundamental limitations of existing approaches*
 - Compare performance and design decisions

Approach	Algorithm	Parallelism			Linearizable	Programming Difficulty
		RR	RW	WW		
Lock	mutex	×	×	×	✓	medium
	rwlock	●	×	×	✓	medium
Lock-free	Harris list [29]	●	●	△	✓	high
Delegation-style	ffwd [55]	×	×	×	✓	low
	NR [11]	●	×	×	✓	low
STM	SwissTM [23]	●	▲	△	✓	lowest
	STO [35]	●	▲	△	✓	low
RCU-style	RCU [46]	●	●	×	-	high
	RLU [44]	●	●	△	-	medium
	MV-RLU	●	●	△	-	medium

NOTE. ×: no parallelism ●: full parallelism ▲: read-write conflict for the same data △: write-write conflict for the same data V: number of versions at GC

Table 1: High-level comparison of synchronization mechanisms. Each mechanism has a unique design goal, strategy to scale, and target class

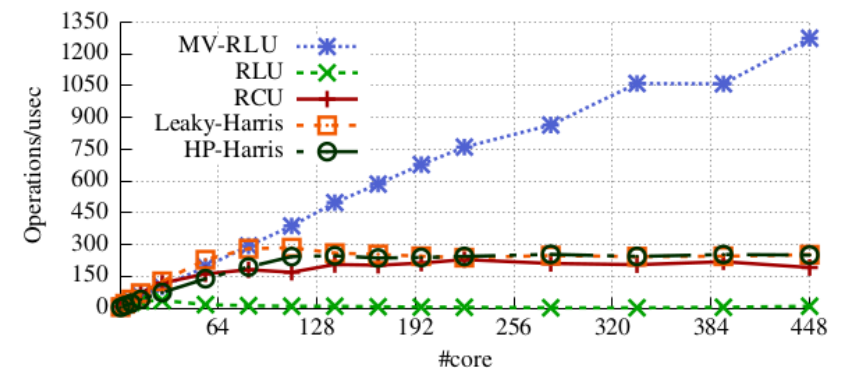


Figure 1: Performance comparison of concurrent hash tables having synchronous log writing (rlu_synchronize) and version chain traversal

Related Work

- ***Show you considered all related work***
- Categorize related papers in a few groups
 - E.g., 1) In-memory database system, 2) Software transactional memory, 3) Synchronization framework for CDS
- Don't fight against each paper
 - The authors will review your paper. ;-(
- Instead appreciate what you inherit from previous work and emphasize you enhanced
- Be precise
- Double check if there are missing related work
 - Especially, PC member's work

Discussion and Limitation

- There is no such perfect system.
- ***Show you made a informed decision even knowing your limitation.***
 - If you do not tell your limitation, reviewers will make their best efforts to figure out your limitation anyway. ;-)
- Explain why you made such decision even if you knew limitations

7. Discussion and Limitations

One of the limitations of MV-RLU is its weaker consistency guarantee: snapshot isolation, which might restrict its use in some applications that requires a stronger consistency guarantee, such as serializability or linearizability. However, the profound adoption of RCU in the Linux kernel, and several database systems, such as Oracle [3], have shown that linearizability is not a necessity. Further, the aforementioned examples also corroborates our view that snapshot isolation is a practical choice for better performance than having stricter consistency

Introduction

- How important it is
 - First 15 minutes of a movie
 - First 30 seconds of a pop song
- You SHOULD make your reviewer excited and jealous
 - “I wish I was working on this area”
 - “If I were working in that area, I wish I was working on this problem”
 - “I wish I was him/her”
- Show your enthusiasm
 - Why this is an important problem
 - Why your solution is good
 - Motivating example
- Articulate technical challenges and your contribution
 - Help review writing

Conclusion and Abstract

- Conclusion
 - Focus on what you achieved and its implication in the future
- Abstract
 - Focus on the problem
 - Introduce your cool, key idea

Revising a paper

- Ask a non-expert colleague to read
 - Ask him/her to read for an hour
 - Understand why they could not understand
 - Understand what aspects of your paper they like
- Revise the paper
- Ask an expert colleague to read
 - Ask him/her to read for 2-3 hours
 - Ask which parts are misleading
 - Understand why they could not understand
 - Understand what aspects of your paper they like
- Do not ask check grammar or spelling

Before submitting a paper

- Title! Title! Title!
 - Reviewer will choose paper only seeing the paper title
 - Misleading title → wrong reviewer assignment → bad review
- Check if every is PERFECT
 - Spelling, grammar, graphic, citations
 - Sloppiness undermines reader's trust
- Revise abstract, introduction, and conclusion

If your paper is rejected

- Paper rejection is a part of research process
 - You just got feedback from research community
- Don't blame reviewers. It's your fault.
 - If they didn't understand, it was your fault for them to misunderstand.
 - If they found flaws of your system, you should thank.
 - Check if reviewers read the paper as you intended
- Finish your homework before next submission
 - Top conferences NEVER accept a paper having clear flaws.

Final thoughts

- Ask a bold question and find a simple solution
- People remember only what you finish
- Find colleagues and mentors
- Help each other
- Keep trying
- Enjoy!